

Loop Optimizations

Nikhil Hegde

Compiler Optimizations course @ QUALCOMM India Pvt. Ltd.

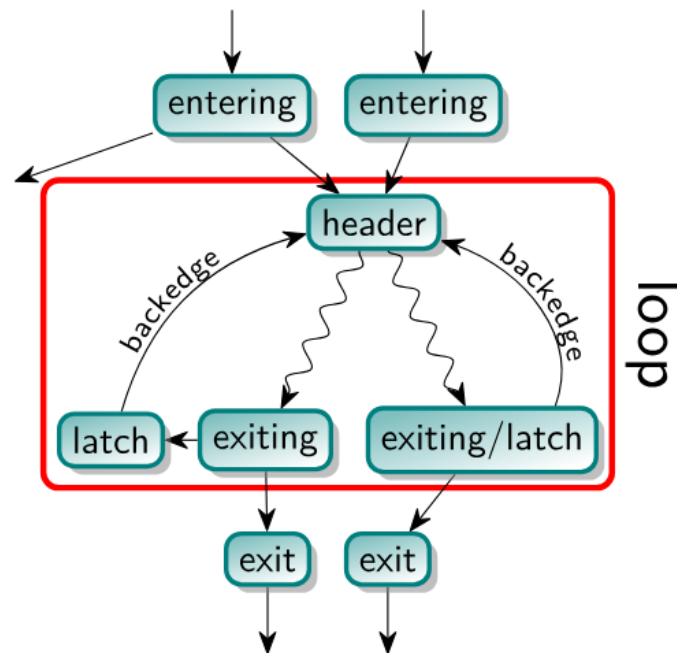
Loop optimization Techniques

- Compiler based
 - Tiling, Fusion, Distribution, LICM
 - Polyhedral compilation
 - Source-to-source transformation: PLUTO, ROSE
- Library based
 - MKL, GotoBLAS
- DSL
 - Halide, SQL, etc.

Recap: Natural loops

- There is a cycle in the CFG (necessary but not sufficient)
- There is only one vertex with in-edges from outside the set. That one vertex is called the **loop-entry/loop-header** block. (also the set needs to be **strongly connected** i.e. there is a path from every vertex to every other vertex in the set).
- For a *backedge* $A \rightarrow B$, the *smallest* set of vertices L including A and B such that $\forall v \in L, predecessors(v) \subseteq L$ or $v = B$
 - **Backedge** – an edge from A to B , where B dominates A
 - How to find backedges?
Do a breadth-first traversal of CFG. Edges that you encounter when discovering new vertices are forward edges. Everything else are backedges.

Natural loops – LLVM Terminology

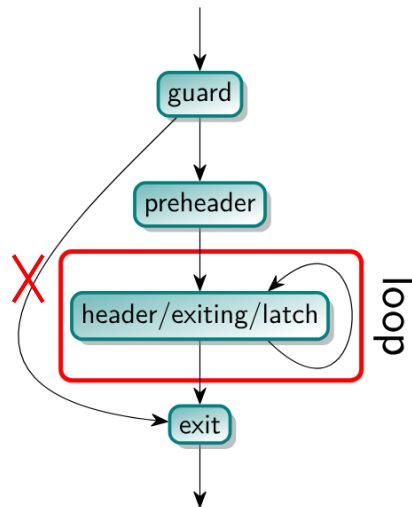


source: <https://llvm.org/docs/LoopTerminology.html>

- Entering block – BB, that has an edge to a loop node (header). Also called preheader if there is only one such block.
- Latch – loop node that has the edge to header (/ source of a backedge)
- Backedge – edge from latch to header
- Exiting block – BB that takes you to a node, which is outside the loop
- Exit block – BB that

Canonical form of loops

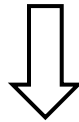
- Adheres to certain restrictions that make it simpler to analyse loops
- E.g. loop simplify form (`passes=loop-simplify`)
 - Has one preheader
 - Has one latch
 - All loop exits (exit blocks) are dominated by loop header



Canonical form of loops

- E.g. loop rotation
 - Converts to a do-while style loop

```
void test(int n) {  
    for (int i = 0; i < n; i += 1)  
        // Loop body  
}
```



```
void test(int n) {  
    int i = 0;  
    do {  
        // Loop body  
        i += 1;  
    } while (i < n);  
}
```

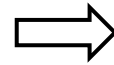
Try it yourself

- For a sample program:
 1. Use the indvars pass to simplify induction variables
 2. Use the loop-simplify pass to view loops in canonical form

Loop Optimizations

- Loop Invariant Code Motion (passes=licm)
 - Saw this earlier while looking at an application of computing dominator info

```
for I = 1 to 100
  for J = 1 to 100
    for K = 1 to 100
      A[I][J][K] = (I*J)*K
```



```
for I = 1 to 100
  temp3=A[I]
  for J = 1 to 100
    temp1=temp3[J]
    temp2=I*J
    for K = 1 to 100
      temp1[K] = temp2*K
```

Certain conditions must hold before factoring out invariant expressions (refer to session_2 slides)

Loop Optimizations

- Induction Variable Simplification (passes=indvars)

```
for(i=0;i<100;i++)  
    foo(a[i])  
      
    ⇒  
    a_end=a+stride*100  
    for(a_i=a;a_i<a_end;a_i+=stride)  
        foo(a_i)
```

- stride is the size of array element.
- `a[i] = a + i*stride` //involves an addition and multiplication
- Transformation eliminates costly multiplication
- Demo: make `indvarssimplify` in `codeexamples`
- Usually followed by Dead Code Elimination and Copy Propagation

Loop Optimizations

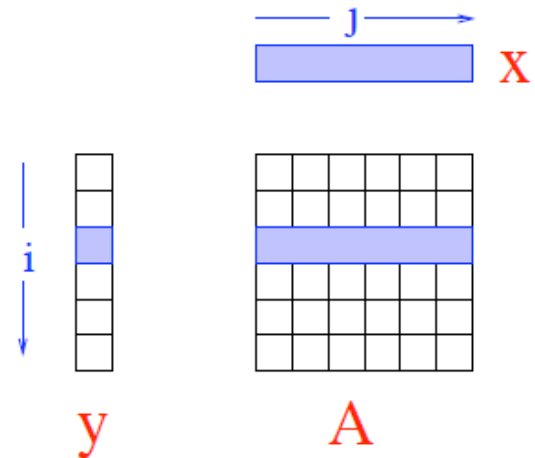
- Low level optimization
 - Moving code around in a single loop
 - Examples: loop invariant code motion, strength reduction, loop unrolling
- High level optimization
 - Restructuring loops, often affects multiple loops
 - Examples: loop fusion, loop interchange, loop tiling

High level loop optimizations

- Many useful compiler optimizations require *restructuring* loops or sets of loops
 - Combining two loops together (*loop fusion*)
 - Switching the order of a nested loop (*loop interchange*)
 - Completely changing the traversal order of a loop (*loop tiling*)

Cache behavior

- Most loop transformations target cache performance
 - Attempt to increase *spatial* or *temporal* locality
 - Locality can be exploited when there is reuse of data (for temporal locality) or recent access of nearby data (for spatial locality)
- Loops are a good opportunity for this: many loops iterate through matrices or arrays
- Consider matrix-vector multiply example
 - Multiple traversals of vector: opportunity for spatial and temporal locality
 - Regular access to array: opportunity for spatial locality



$$y = Ax$$

```
for (i = 0; i < N; i++)  
    for (j = 0; j < N; j++)  
        y[i] += A[i][j] * x[j]
```

Loop fusion

```
do i = 1, n
```

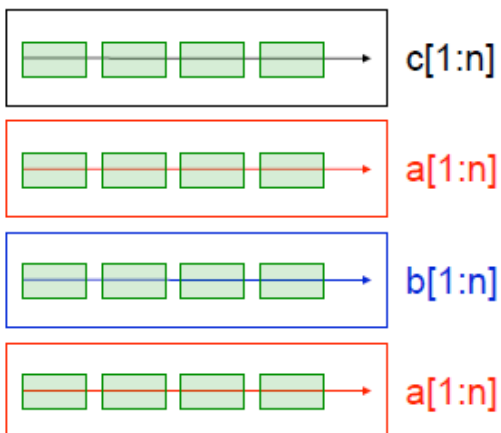
```
  c[i] = a[i]
```

```
end do
```

```
do i = 1, n
```

```
  b[i] = a[i]
```

```
end do
```



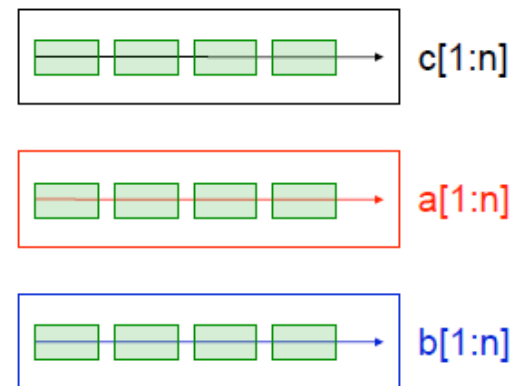
- Combine two loops together into a single loop
- Why is this useful?
- Is this always legal?

```
do i = 1, n
```

```
  c[i] = a[i]
```

```
  b[i] = a[i]
```

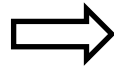
```
end do
```



Loop Fusion – Another Example

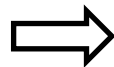
```
for(i=0;i<100;i++)  
    b[i]=foo(a[i])
```

```
for(i=0;i<100;i++)  
    c[i]=bar(b[i])
```



```
for(i=0;i<100;i++)  
    b[i]=foo(a[i])  
    c[i]=bar(b[i])
```

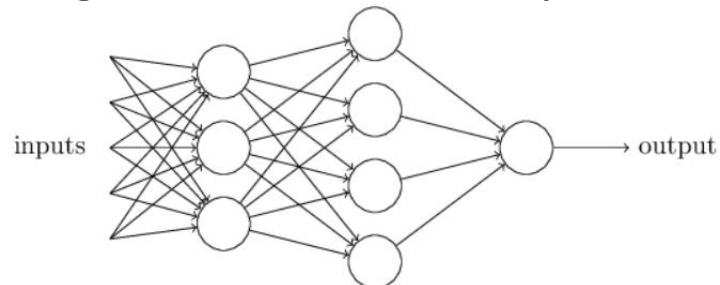
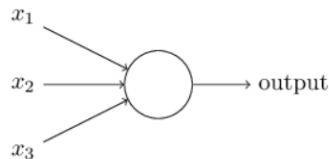
If b is dead outside the loop, then:



```
for(i=0;i<100;i++)  
    c[i]=bar(foo(a[i]))
```

This is a big saving – eliminate memory store, reduce memory footprint etc.

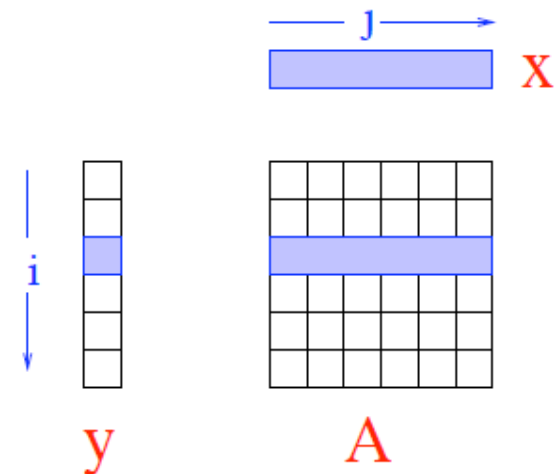
Common in DNN applications e.g. MatMul followed by ReLU



output = weighted sum of inputs, Multiple inputs => Matrix of inputs * Matrix of weights

Loop interchange

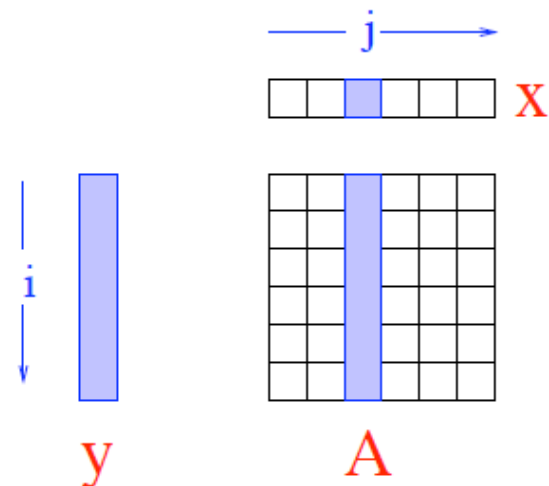
- Change the order of a nested loop
- This is not always legal – it changes the order that elements are accessed!
- Why is this useful?
- Consider matrix-matrix multiply when A is stored in column-major order (i.e., each column is stored in contiguous memory)



```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    y[i] += A[i][j] * x[j]
```

Loop interchange

- Change the order of a nested loop
- This is not always legal – it changes the order that elements are accessed!
- Why is this useful?
- Consider matrix-matrix multiply when A is stored in column-major order (i.e., each column is stored in contiguous memory)



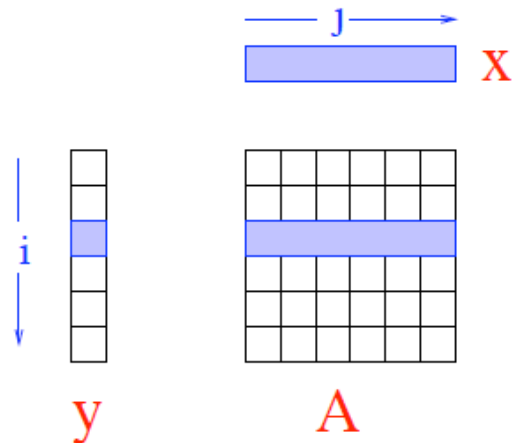
```
for (j = 0; j < N; j++)  
  for (i = 0; i < N; i++)  
    y[i] += A[i][j] * x[j]
```


Loop tiling

- Also called “loop blocking”
- One of the more complex loop transformations
- Goal: break loop up into smaller pieces to get spatial and temporal locality
- Create new inner loops so that data accessed in inner loops fit in cache
- Also changes iteration order, so may not be legal

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    y[i] += A[i][j] * x[j]
```

```
for (ii = 0; ii < N; ii += B)  
  for (jj = 0; jj < N; jj += B)  
    for (i = ii; i < ii+B; i++)  
      for (j = jj; j < jj+B; j++)  
        y[i] += A[i][j] * x[j]
```

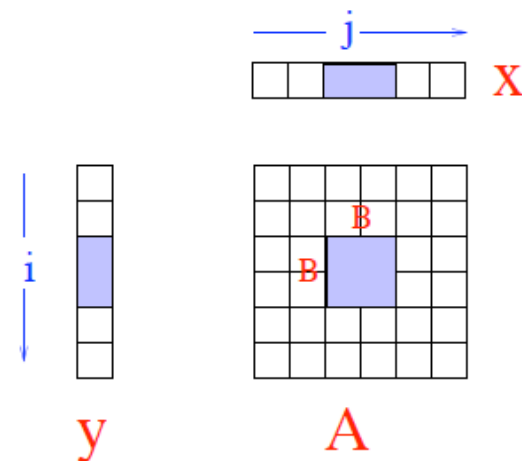


Loop tiling

- Also called “loop blocking”
- One of the more complex loop transformations
- Goal: break loop up into smaller pieces to get spatial and temporal locality
- Create new inner loops so that data accessed in inner loops fit in cache
- Also changes iteration order, so may not be legal

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    y[i] += A[i][j] * x[j]
```

```
for (ii = 0; ii < N; ii += B)  
  for (jj = 0; jj < N; jj += B)  
    for (i = ii; i < ii+B; i++)  
      for (j = jj; j < jj+B; j++)  
        y[i] += A[i][j] * x[j]
```



Legality of High Level Loop Optimizations

- Dependence Analysis (slide courtesy: Milind Kukarni)

Dependence Analysis

Motivating question

- Can the loops on the right be run in parallel?
- i.e., can different processors run different iterations in parallel?
- What needs to be true for a loop to be parallelizable?
- Iterations cannot interfere with each other
- No *dependence* between iterations

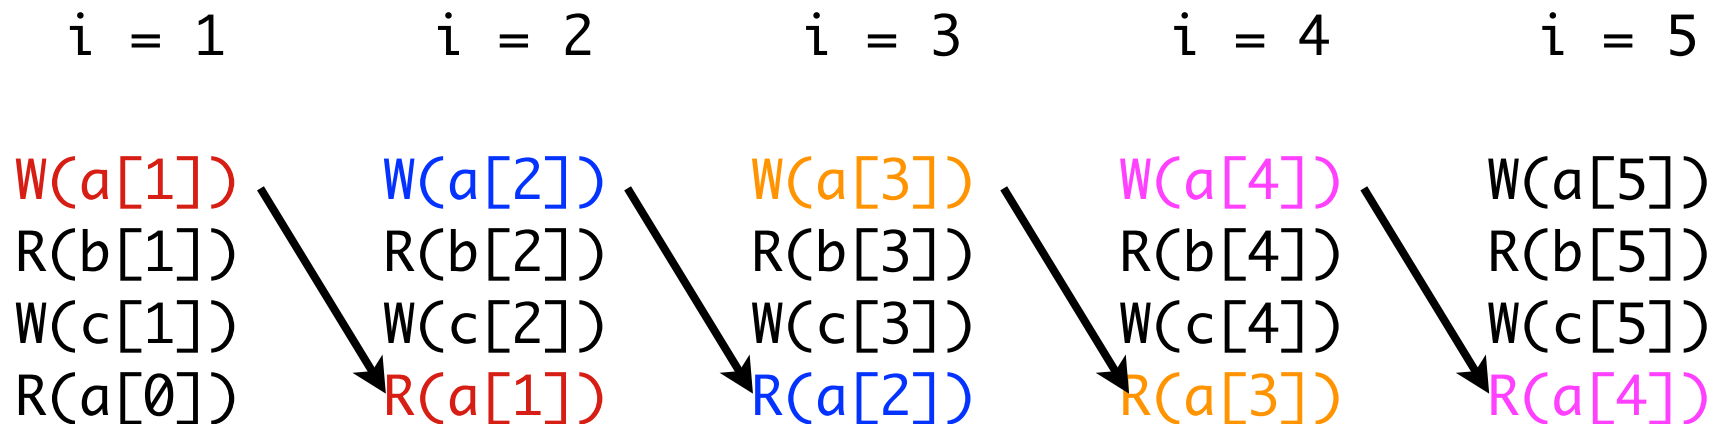
```
for (i = 1; i < N; i++) {  
    a[i] = b[i];  
    c[i] = a[i - 1];  
}
```

```
for (i = 1; i < N; i++) {  
    a[i] = b[i];  
    c[i] = a[i] + b[i - 1];  
}
```

Dependences

- A *flow dependence* occurs when one iteration writes a location that a *later* iteration reads

```
for (i = 1; i < N; i++) {  
    a[i] = b[i];  
    c[i] = a[i - 1];  
}
```



Running a loop in parallel

- If there is a dependence in a loop, we cannot guarantee that the loop will run correctly in parallel
- What if the iterations run out of order?
 - Might read from a location before the correct value was written to it
- What if the iterations do not run in lock-step?
 - Same problem!

Other kinds of dependence

- *Anti dependence* – When an iteration *reads* a location that a later iteration *writes* (why is this a problem?)

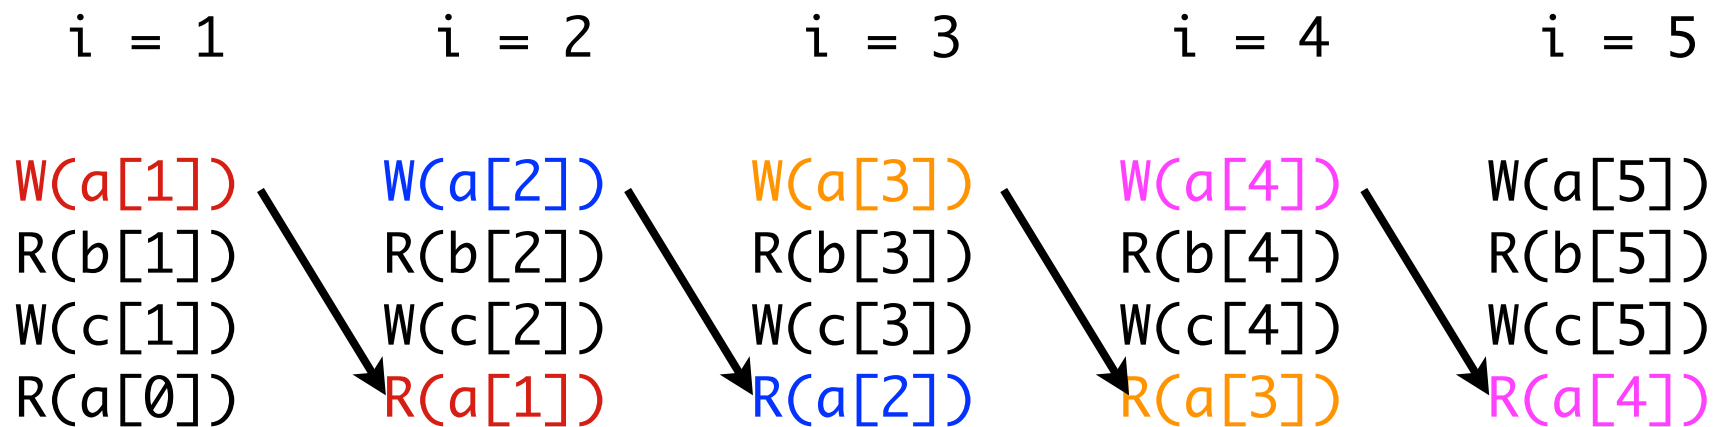
```
for (i = 1; i < N; i++) {  
    a[i - 1] = b[i];  
    c[i] = a[i];  
}
```

- *Output dependence* – When an iteration *writes* a location that a later iteration *writes* (why is this a problem?)

```
for (i = 1; i < N; i++) {  
    a[i] = b[i];  
    a[i + 1] = c[i];  
}
```


Data dependence concepts

- Dependence *source* is the earlier statement (the statement at the tail of the dependence arrow)
- Dependence *sink* is the later statement (the statement at the head of the dependence arrow)



- Dependences can only go forward in time: always from an earlier iteration to a later iteration.

Using dependences

- If there are no dependences, we can parallelize a loop
 - None of the iterations interfere with each other
- Can also use dependence information to drive other optimizations
 - Loop interchange
 - Loop fusion
 - (We will discuss these later)
- Two questions:
 - How do we represent dependences in loops?
 - How do we determine if there are dependences?

Representing dependences

- Focus on flow dependences for now
- Dependences in straight line code are easy to represent:
 - One statement writes a location (variable, array location, etc.) and another reads that same location
 - Can figure this out using reaching definitions
- What do we do about loops?
- We often care about dependences between the same statement in different iterations of the loop!

```
for (i = 1; i < N; i++) {  
    a[i + 1] = a[i] + 2  
}
```

Iteration space graphs

- Represent each *dynamic* instance of a loop as a point in a graph
- Draw arrows from one point to another to represent dependences

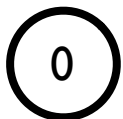
```
for (i = 0; i < N; i++) {  
    a[i + 2] = a[i]  
}
```

Iteration space graphs

- Represent each *dynamic* instance of a loop as a point in a graph
- Draw arrows from one point to another to represent dependences

```
for (i = 0; i < N; i++) {  
    a[i + 2] = a[i]  
}
```

- Step 1: Create nodes, 1 for each iteration
 - Note: not 1 for each array location!

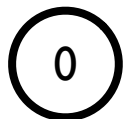


Iteration space graphs

- Represent each *dynamic* instance of a loop as a point in a graph
- Draw arrows from one point to another to represent dependences

```
for (i = 0; i < N; i++) {  
    a[i + 2] = a[i]  
}
```

- Step 2: Determine which array elements are read and written in each iteration



R: a[0]
W: a[2]



R: a[1]
W: a[3]



R: a[2]
W: a[4]



R: a[3]
W: a[5]



R: a[4]
W: a[6]



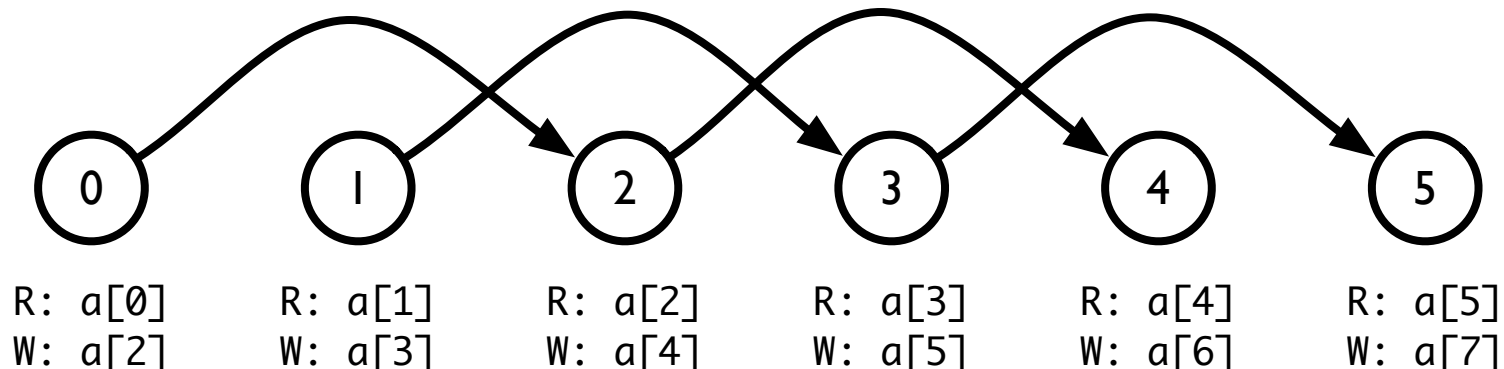
R: a[5]
W: a[7]

Iteration space graphs

- Represent each *dynamic* instance of a loop as a point in a graph
- Draw arrows from one point to another to represent dependences

```
for (i = 0; i < N; i++) {  
    a[i + 2] = a[i]  
}
```

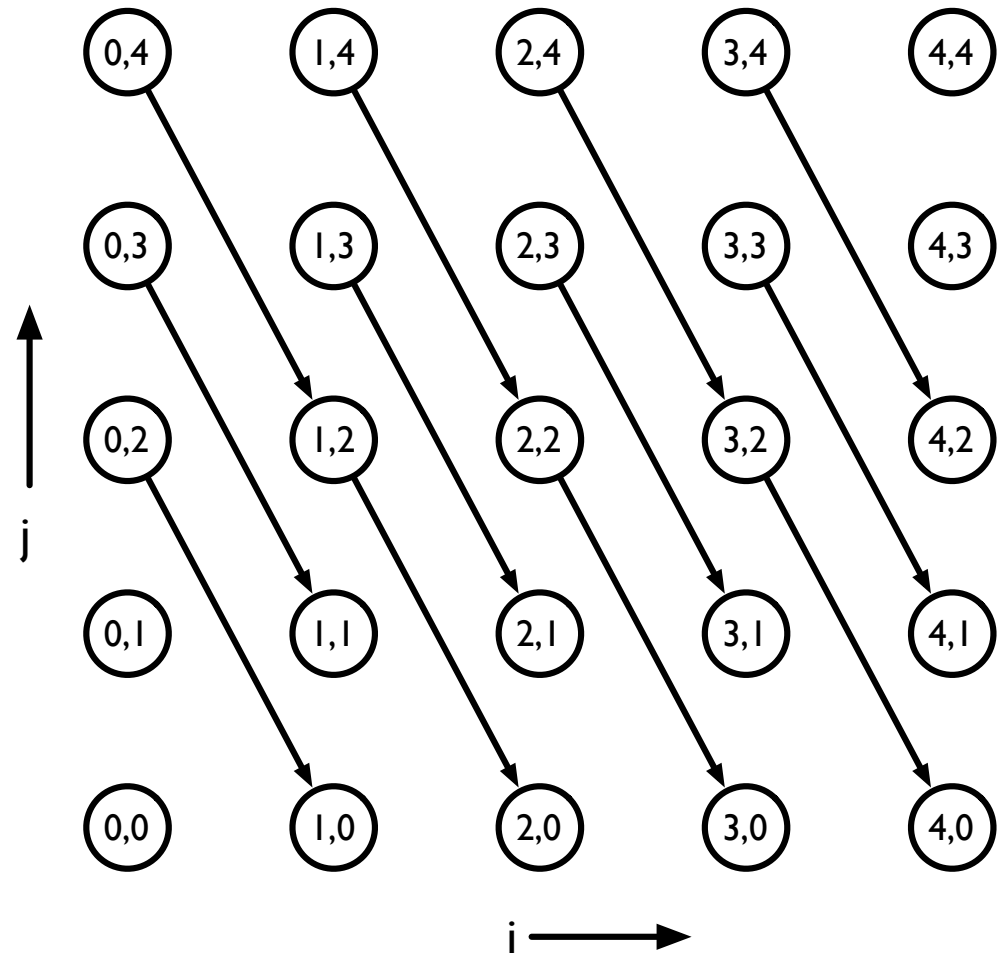
- Step 3: Draw arrows to represent dependences





2-D iteration space graphs

- Can do the same thing for doubly-nested loops
- 2 loop counters

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    a[i+1][j-2] = a[i][j] + 1
```



Iteration space graphs

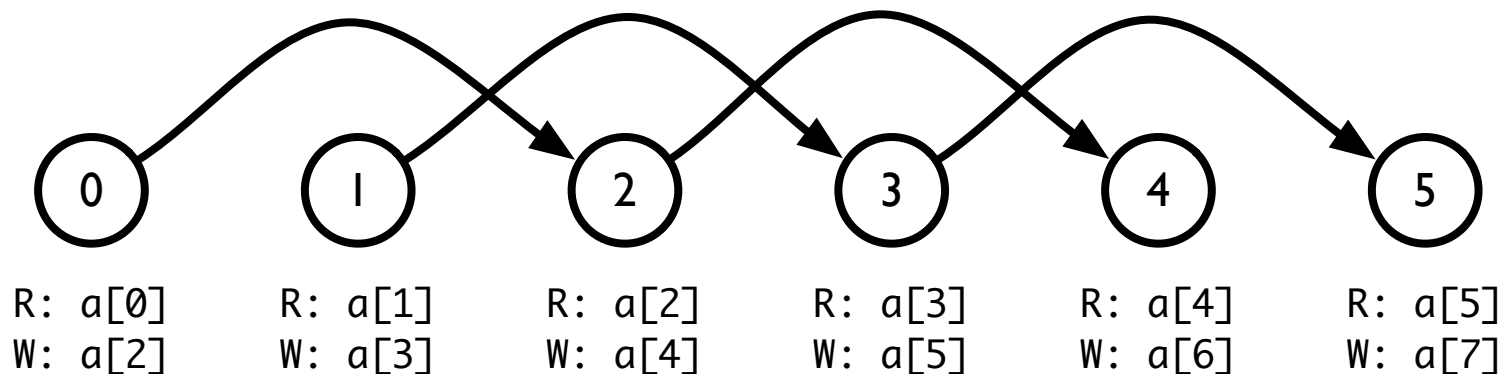
- Can also represent output and anti dependences
 - Use different kinds of arrows for clarity. *E.g.*
 -  for output
 -  for anti
- Crucial problem: Iteration space graphs are potentially infinite representations!
- Can we represent dependences in a more compact way?

Distance and direction vectors

- Compiler researchers have devised *compressed* representations of dependences
 - Capture the same dependences as an iteration space graph
 - May lose *precision* (show more dependences than the loop actually has)
- Two types
 - Distance vectors: captures the “shape” of dependences, but not the particular source and sink
 - Direction vectors: captures the “direction” of dependences, but not the particular shape

Distance vector

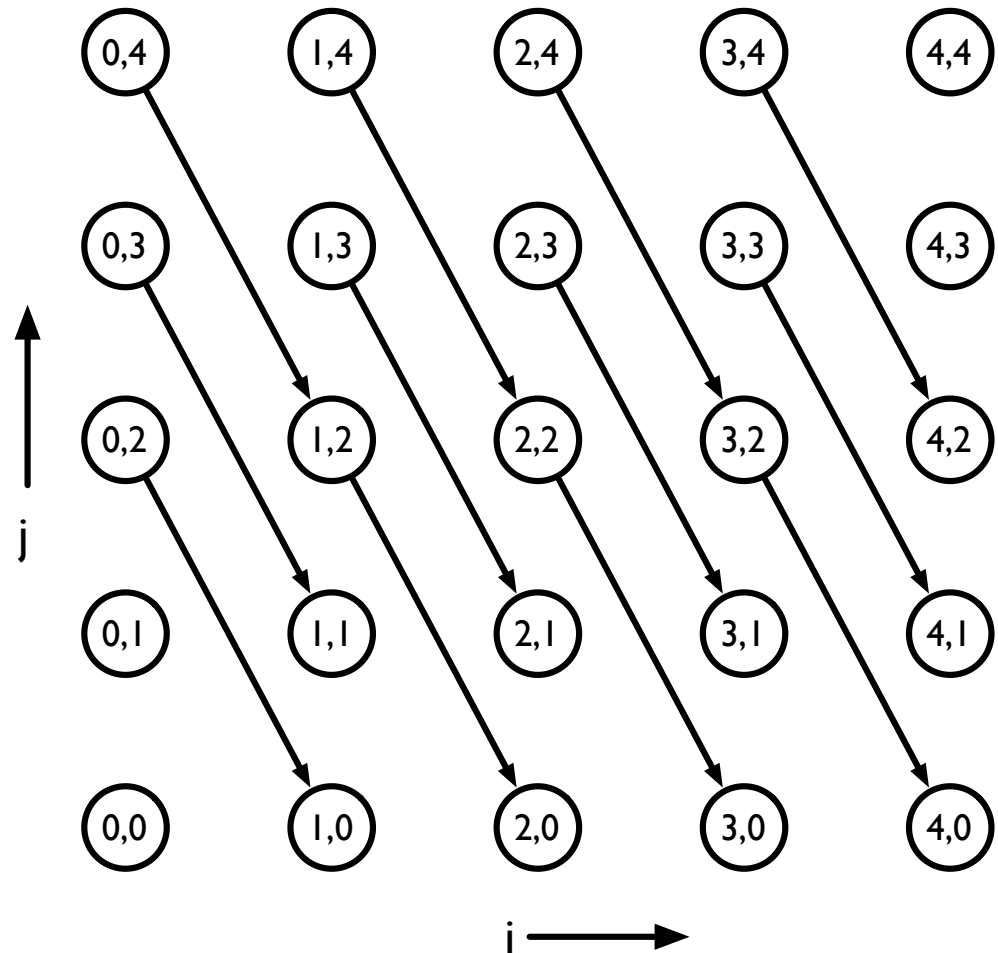
- Represent each dependence arrow in an iteration space graph as a vector
- Captures the “shape” of the dependence, but loses where the dependence originates



- Distance vector for this iteration space: (2)
 - Each dependence is 2 iterations forward

2-D distance vectors

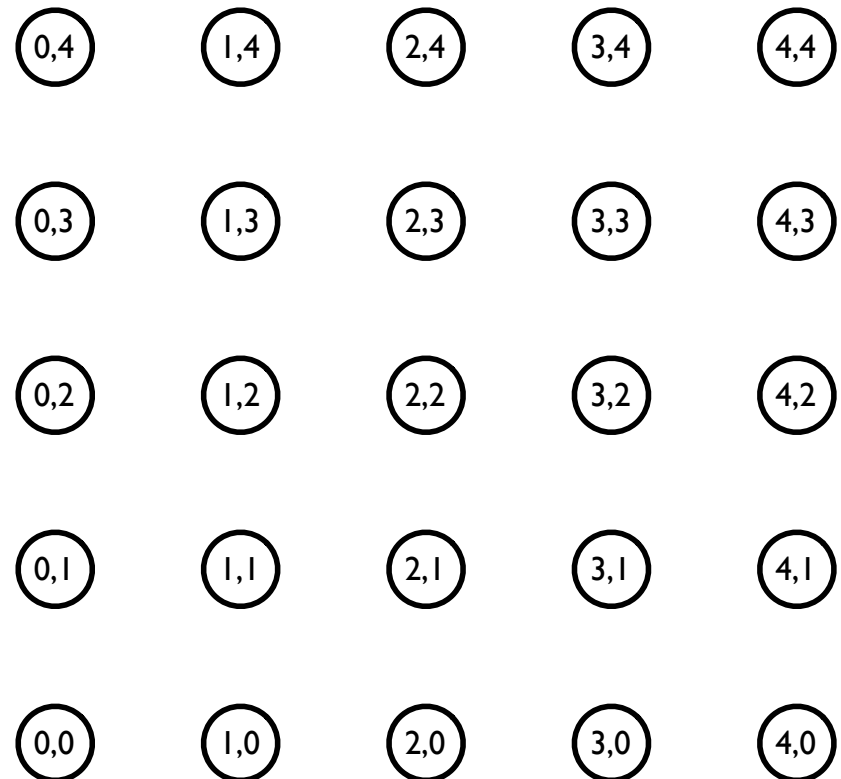
- Distance vector for this graph:
 - $(1, -2)$
 - +1 in the i direction, -2 in the j direction
- Crucial point about distance vectors: they are always “positive”
- First non-zero entry has to be positive
- Dependences can't go backwards in time



More complex example

- Can have multiple distance vectors

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    a[i+1][j-2] = a[i][j] +  
                  a[i-1][j-2]
```

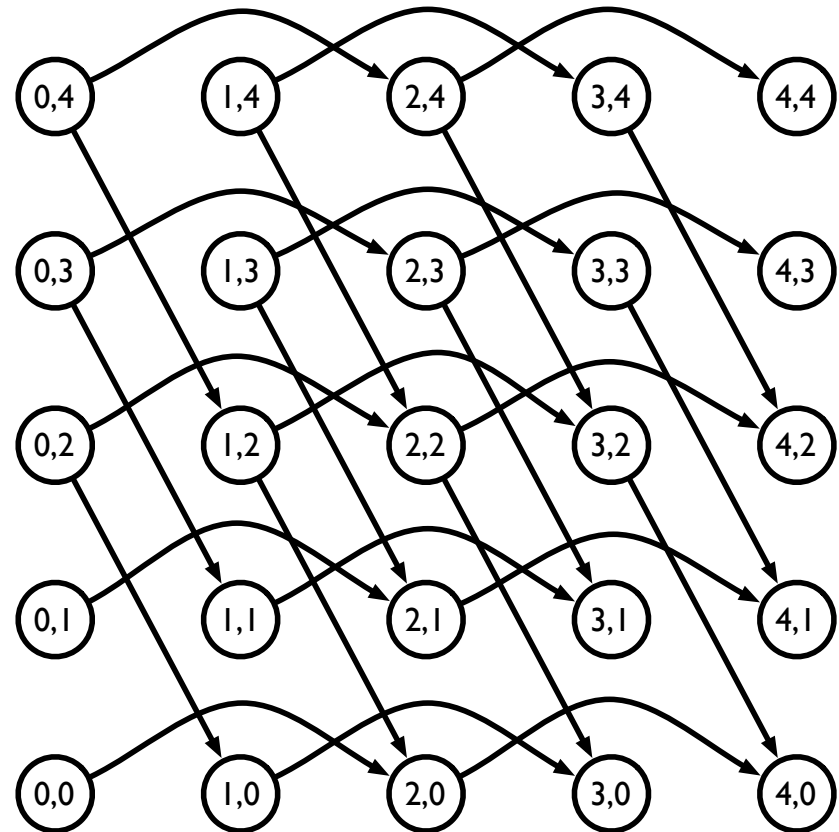


More complex example

- Can have multiple distance vectors

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    a[i+1][j-2] = a[i][j] +  
                  a[i-1][j-2]
```

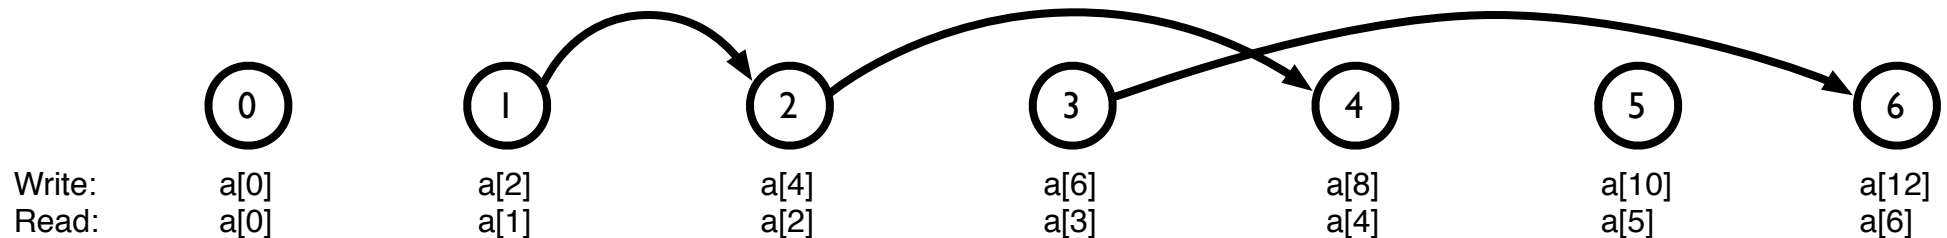
- Distance vectors
 - (1, -2)
 - (2, 0)
- Important point: order of vectors depends on order of loops, not use in arrays



Problems with distance vectors

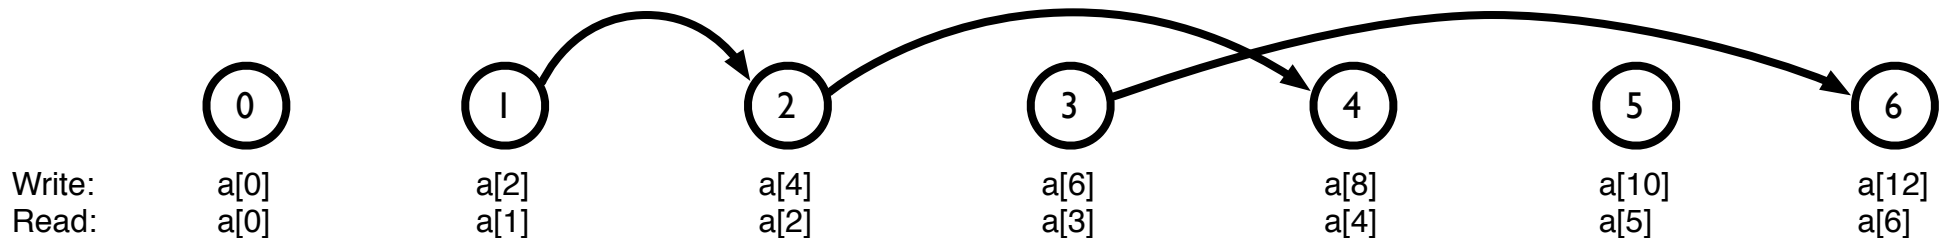
- The preceding examples show how distance vectors can summarize all the dependences in a loop nest using just a small number of distance vectors
- Can't always summarize as easily
- Running example:

```
for (i = 0; i < N; i++)  
    a[2*i] = a[i];
```



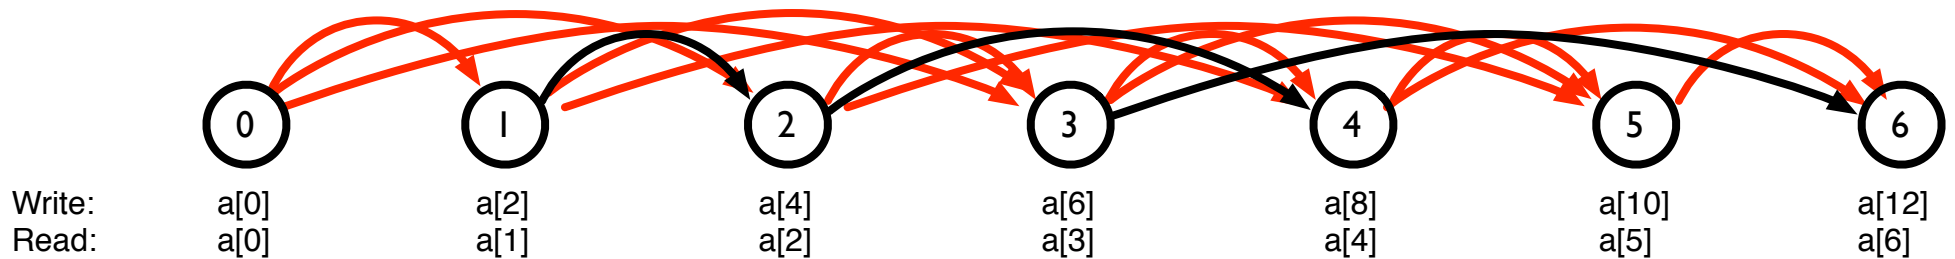
Loss of precision

- What are the distance vectors for this code?
 - (1), (2), (3), (4) ...
- Note: we have information about the length of each vector, but not about the source of each vector
- What happens if we try to reconstruct the iteration space graph?



Loss of precision

- What are the distance vectors for this code?
- (1), (2), (3), (4) ...
- Note: we have information about the length of each vector, but not about the source of each vector
- What happens if we try to reconstruct the iteration space graph?



Direction vectors

- The whole point of distance vectors is that we want to be able to succinctly capture the dependences in a loop nest
 - But in the previous example, not only did we add a lot of extra information, we still had an infinite number of distance vectors
- Idea: summarize distance vectors, and save only the *direction* the dependence was in
 - $(2, -1) \rightarrow (+, -)$
 - $(0, 1) \rightarrow (0, +)$
 - $(0, -2) \rightarrow (0, -)$
 - (can't happen; dependences have to be positive)
- Notation: sometimes use ' $<$ ' and ' $>$ ' instead of '+' and '-'

Why use direction vectors?

- Direction vectors lose a lot of information, but do capture some useful information
 - Whether there is a dependence (anything other than a '0' means there is a dependence)
 - Which dimension and direction the dependence is in
- Many times, the only information we need to determine if an optimization is legal is captured by direction vectors
 - Loop parallelization
 - Loop interchange

Loop parallelization

Loop-carried dependence

- The key concept for parallelization is the *loop carried dependence*
- A dependence that crosses loop iterations
- If there is a loop carried dependence, then that loop *cannot* be parallelized
- Some iterations of the loop depend on other iterations of the same loop

Examples

```
for (i = 0; i < N; i++)  
    a[2*i] = a[i];
```

Later iterations of i loop
depend on earlier iterations

```
for (i = 0; i < N; i++)  
    for (j = 0; j < N; j++)  
        a[i+1][j-2] = a[i][j] + 1
```

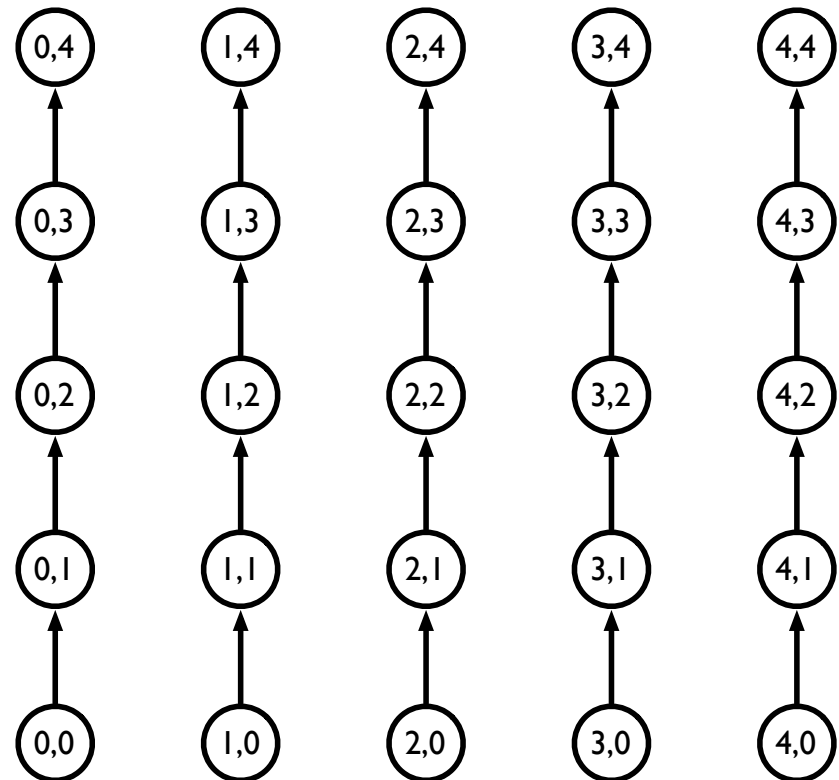
Later iterations of both i and
j loops depend on earlier iterations

Some subtleties

- Dependences might only be carried over one loop!

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    a[i][j+1] = a[i][j] + 1
```

- Can parallelize i loop, but not j loop

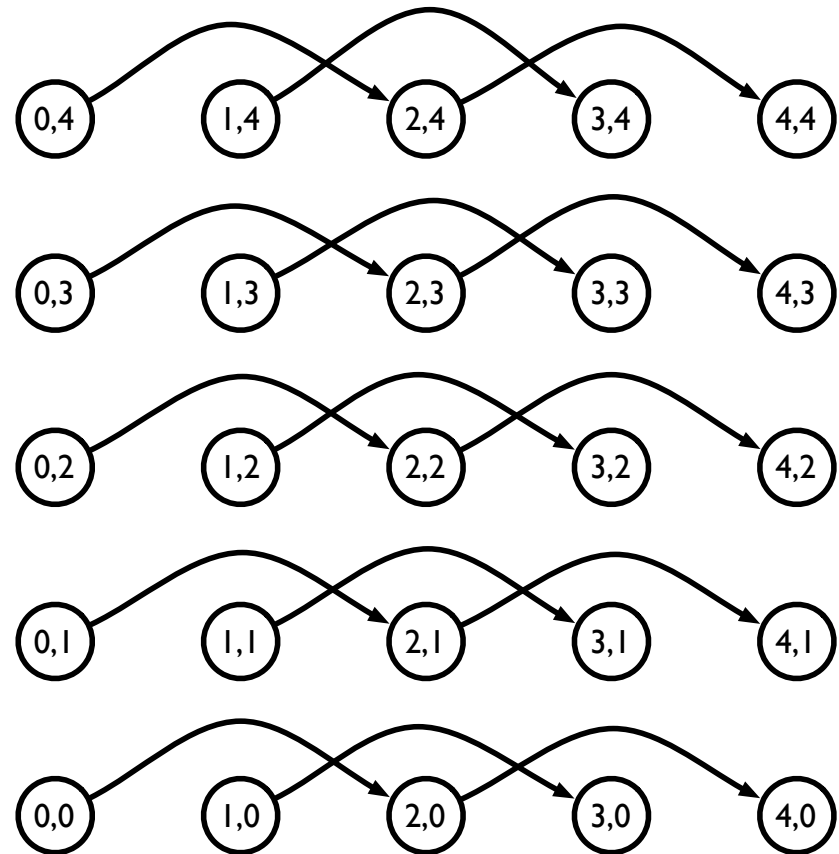


Some subtleties

- Dependences might only be carried over one loop!

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    a[i+1][j] = a[i-1][j] + 1
```

- Can parallelize j loop, but not i loop



Direction vectors

- So how do direction vectors help?
 - If there is a non-zero entry for a loop dimension, that means that there is a loop carried dependence over that dimension
 - If an entry is zero, then that loop can be parallelized!
- May be able to parallelize inner loop even if entry is not zero, but you have to carefully structure parallel execution

Other loop optimizations

Loop interchange

- We've seen this one before
- Interchange doubly-nested loop to
 - Improve locality
 - Improve parallelism
 - Move parallel loop to outer loop (coarse grained parallelism)

Loop interchange legality

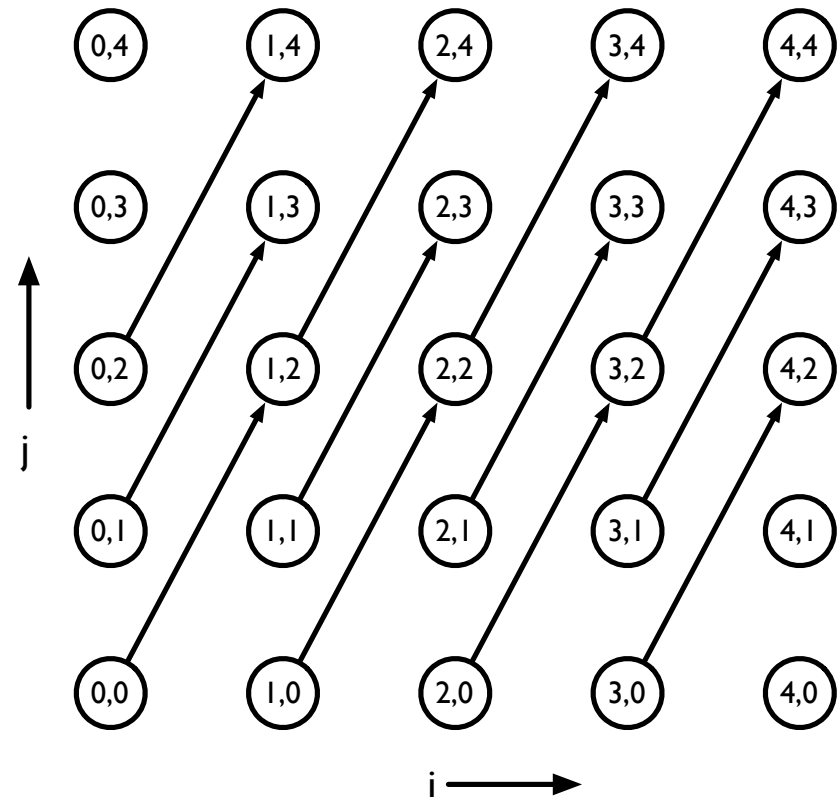
- We noted that loop interchange is not always legal, because it reorders a computation
- Can we use dependences to determine legality?

Loop interchange dependences

- Consider interchanging the following loop, with the dependence graph to the right:

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    a[i+1][j+2] = a[i][j] + 1
```

- Distance vector (1, 2)
- Direction vector (+, +)

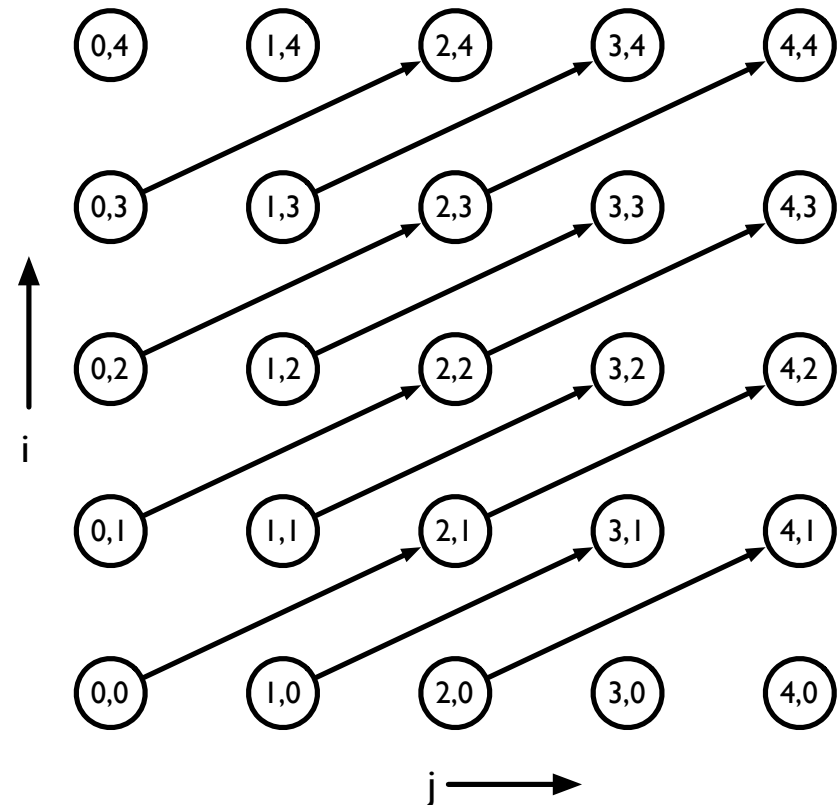


Loop interchange dependences

- Consider interchanging the following loop, with the dependence graph to the right:

```
for (j = 0; j < N; j++)  
  for (i = 0; i < N; i++)  
    a[i+1][j+2] = a[i][j] + 1
```

- Distance vector (2, 1)
- Direction vector (+, +)
- Distance vector gets swapped!



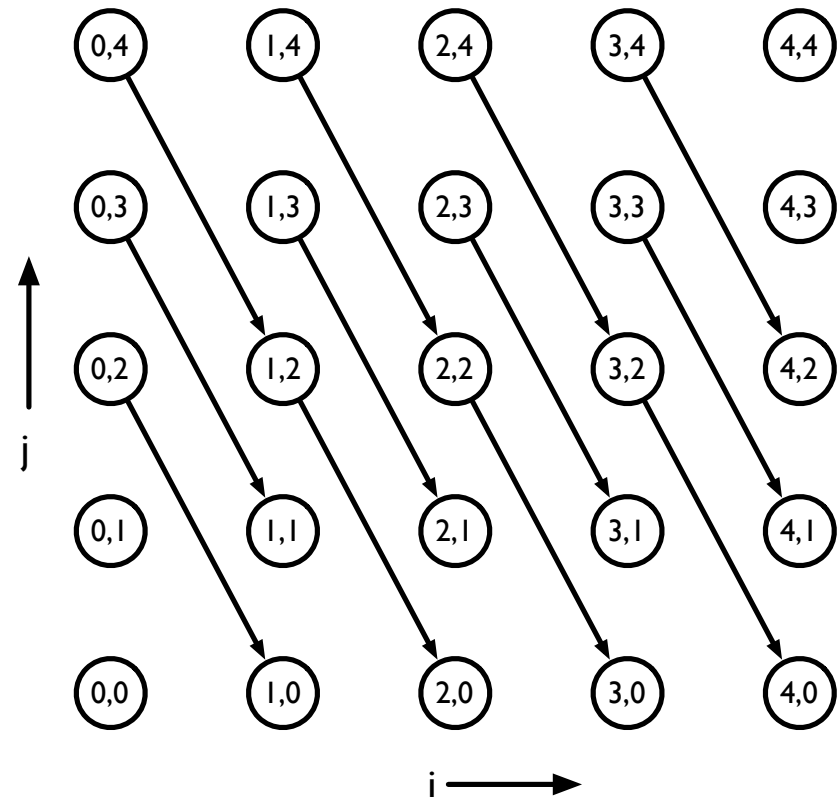
Loop interchange legality

- Interchanging two loops swaps the order of their entries in distance/direction vectors
 - $(0, +) \rightarrow (+, 0)$
 - $(+, 0) \rightarrow (0, +)$
- But remember, we can't have backwards dependences
 - $(+, -) \rightarrow (-, +)$
 - Illegal dependence \rightarrow Loop interchange not legal!

Loop interchange dependences

- Example of illegal interchange:

```
for (i = 0; i < N; i++)  
  for (j = 0; j < N; j++)  
    a[i+1][j-2] = a[i][j] + 1
```

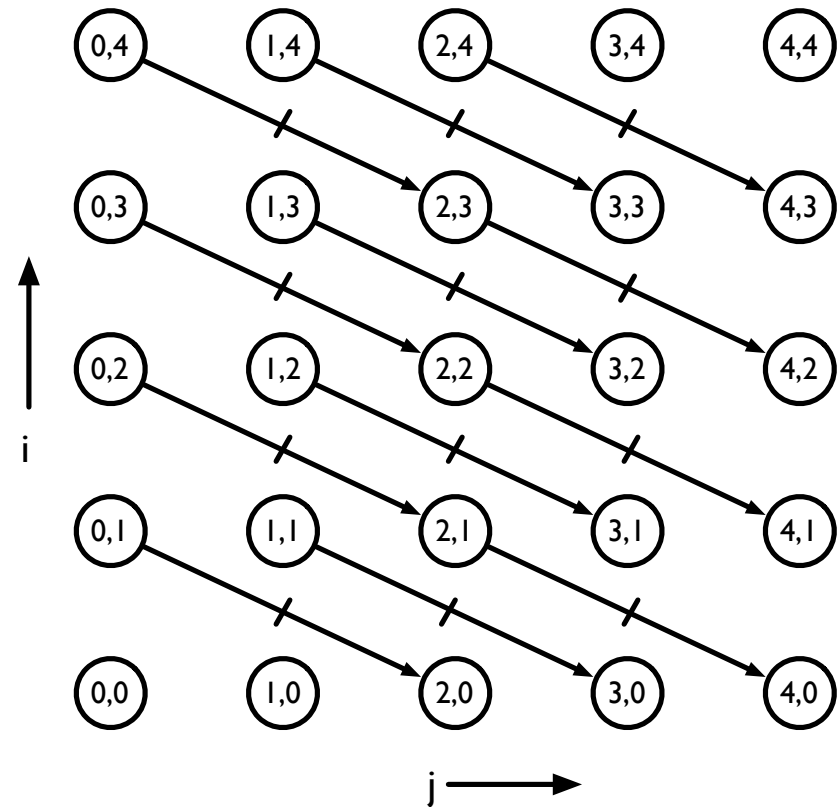


Loop interchange dependences

- Example of illegal interchange:

```
for (j = 0; j < N; j++)  
  for (i = 0; i < N; i++)  
    a[i+1][j-2] = a[i][j] + 1
```

- Flow dependences turned into anti-dependences
- Result of computation will change!



Loop fusion/distribution

- Loop fusion: combining two loops into a single loop
 - Improves locality, parallelism
- Loop distribution: splitting a single loop into two loops
 - Can increase parallelism (turn a non-parallelizable loop into a parallelizable loop)
- Legal as long as optimization maintains dependences
 - Every dependence in the original loop should have a dependence in the optimized loop
 - Optimized loop should not introduce new dependences

Fusion/distribution example

- Code 1:

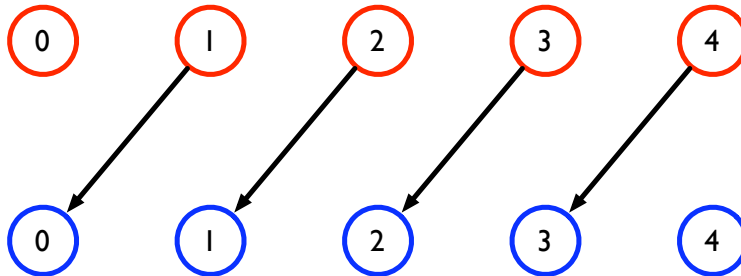
```
for (i = 0; i < N; i++)
```

```
  a[i - 1] = b[i]
```

```
for (j = 0; j < N; j++)
```

```
  c[j] = a[j]
```

- Dependence graph



- All red iterations finish before blue iterations → flow dependence

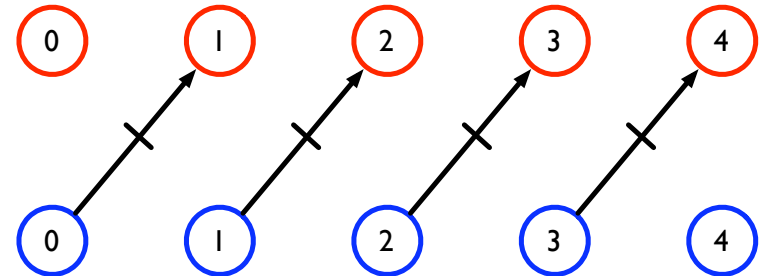
- Code 2:

```
for (i = 0; i < N; i++)
```

```
  a[i - 1] = b[i]
```

```
  c[i] = a[i]
```

- Dependence graph



- i iterations finish before i+1 iterations → flow dependence now an anti dependence!

Fusion/distribution utility

for (i = 0; i < N; i++)
 a[i] = a[i - 1]

→ Fusion

for (i = 0; i < N; i++)
 a[i] = a[i - 1]

← Distribution

for (j = 0; j < N; j++)
 b[j] = a[j]

- Fusion and distribution both legal
- Right code has better locality, but cannot be parallelized due to loop carried dependences
- Left code has worse locality, but blue loop can be parallelized